



Norwegian University of
Science and Technology

Programming Models for Multi-core Systems

Mujahed Eleyat

October 1, 2014

Learning Goals

- Define of a parallel programming model
- Classify parallel programming models
- Explain data parallelism and task parallelism
- Explain several approaches for parallel programming
- Present examples of parallel programming models

The Sequential Execution Model

- Programming was easy
- Applications gain more performance as the frequency of uniprocessors are increased
- The shift to multicore systems put an end to the “free lunch”
 - Automatic “good” utilization of multicore processing power is very difficult



Challenges of Multicore Programming

- Parallel algorithms are hard to develop, implement, and debug
- Efficiency is tied to knowledge architectural details
- Wide range of architectural differences
- Memory Wall and Power Wall
- Fast evolution
- Many tools by many vendors



A Parallel Programming Model

- An abstraction of a computer system
- Specifies the programmer's view on parallel computer
- Influenced by the architectural design and the language, compiler, or the runtime libraries
 - different parallel programming models for the same architecture
- Implementation of a Parallel Programming Model
 - Libraries
 - Compiler directives
 - Parallel languages

Characteristics of an Ideal Multicore Programming Model

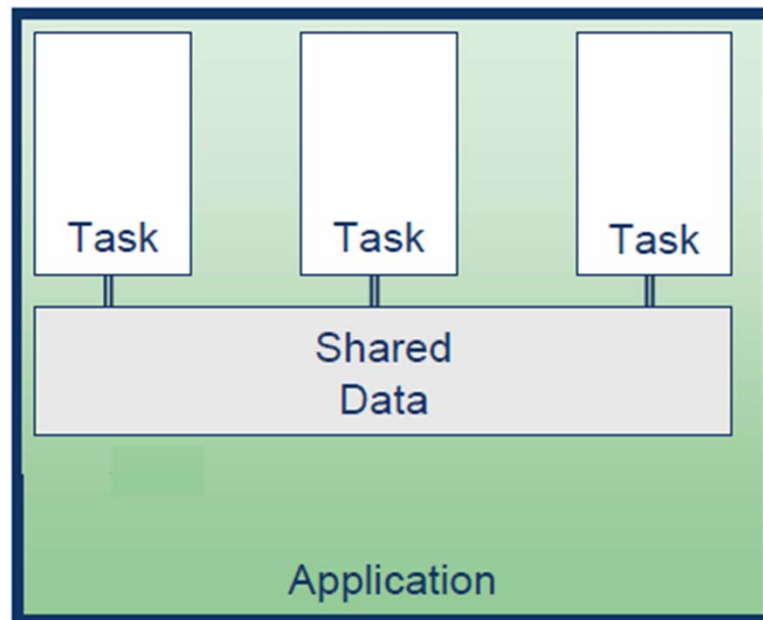
1. Productivity
 - Ease of programming
2. Performance
3. Portability
4. Forward Scalability
 - a) Number of compute resource
 - b) Diversity of computer resource types

Programming Models as Abstractions of Multicore Systems


1. Classical parallel programming models
 - Shared Memory
 - Pthreads and OpenMP
 - Message passing
 - MPI
2. Heterogeneous Parallel Programming Models
3. Partitioned Global Address Space (PGAS)
4. Hybrid, shared-distributed memory + GPUs

Shared Memory and Message Passing

Shared Memory



Message Passing




OpenCL

- An open standard heterogeneous programming model
- Main goals are portability and code reusability across heterogeneous platforms
- Host(CPU) and devices (GPUs, CPUs, ...)
- Robust API for doing data parallel and task parallel computation
- Not easy to use

OpenCL Platform Model

- A host coordinates execution, transferring data to and from an array of Compute Devices.
- Each Compute Device is composed of an array of Compute Units.
- Each Compute Unit is composed of an array of Processing Elements.



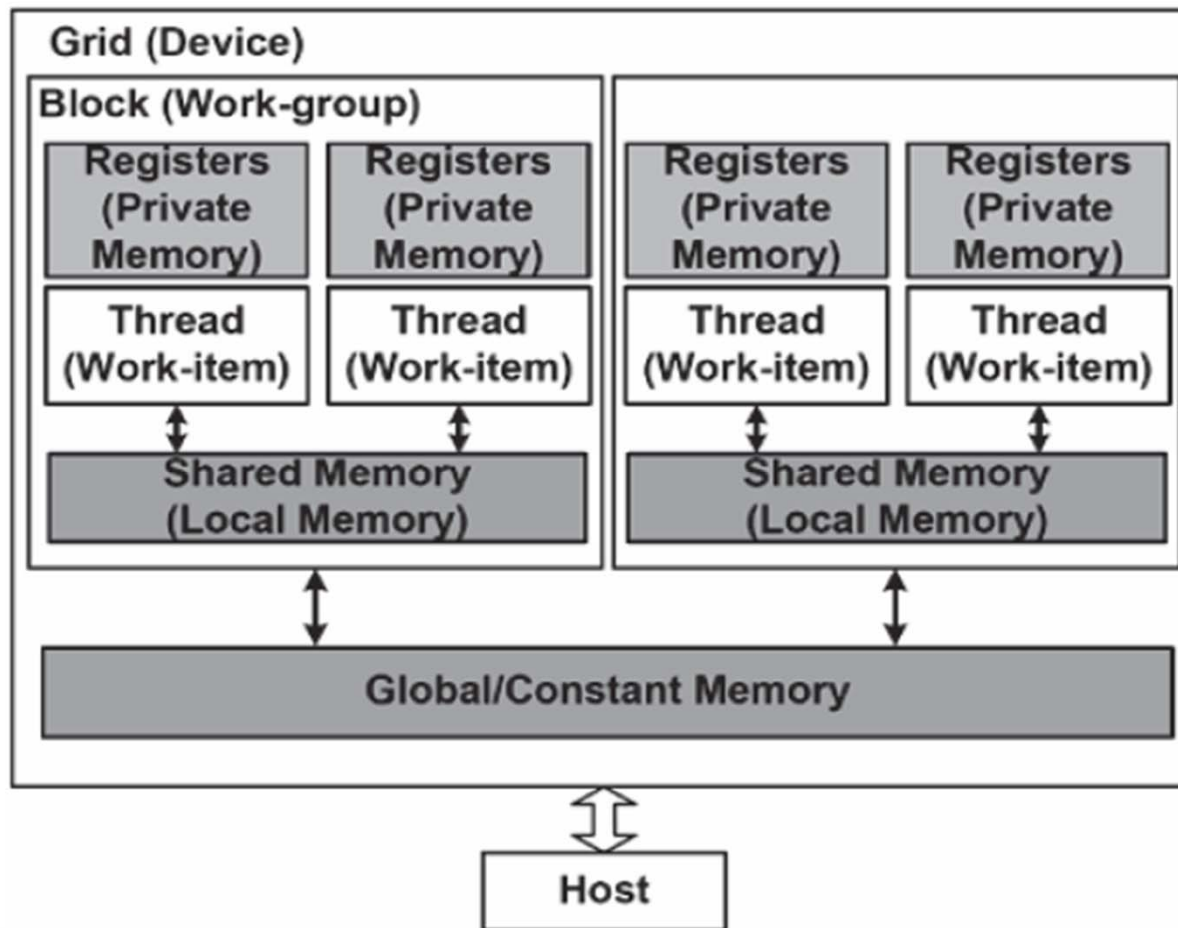
NTNU

Norwegian University of
Science and Technology

OpenCL Execution Model

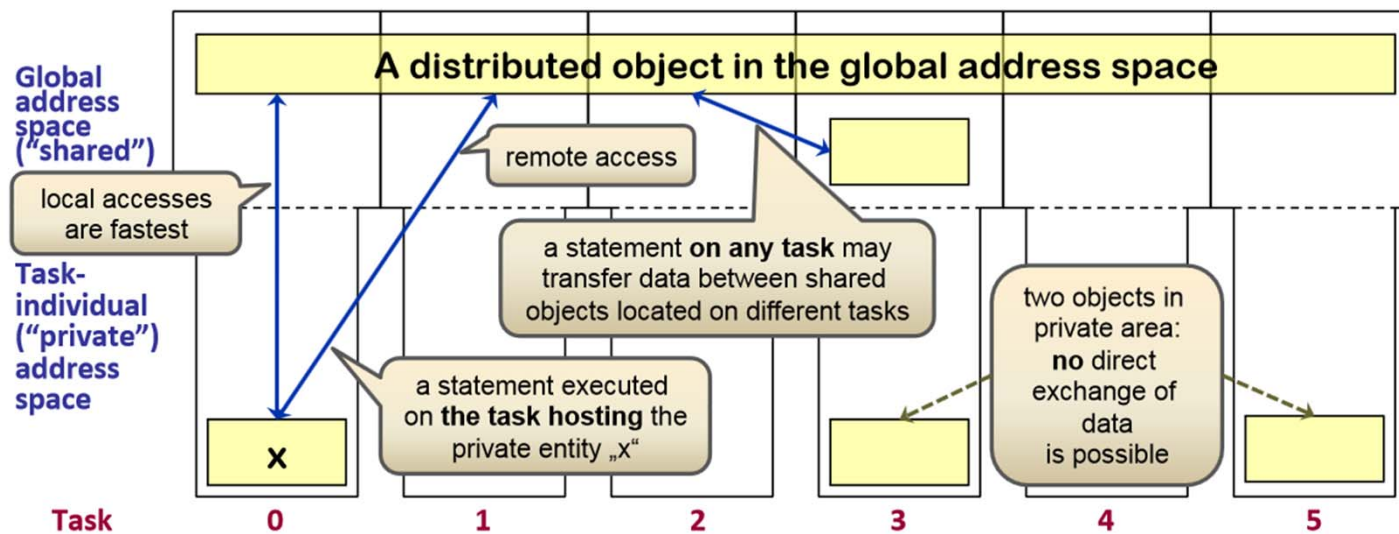
- Simultaneous execution of multiple instances of a kernel on one or more OpenCL devices
- Kernels/(work items) are organized into work groups for communication and synchronization
- The host application enqueues commands to a command queue
 - Kernel execution commands launch work items
- kernels are compiled for the targeted device in runtime
 - The host application exploit all compute devices on the system

OpenCL Memory Model



Partitioned Global Address Space (PGAS)

- Global memory address space that is logically partitioned and a portion of it is local to each process or thread
 - exploiting locality of reference.
- Basis for Unified Parallel C, Coarray Fortran, Fortress, Chapel, X10



Hybrid Parallel Programming Models

- Mixing different programming models as a match to the hybrid hardware architectures
 - Pthreads and MPI (SMP Cluster)
 - MPI and OpenMP (SMP Cluster)
 - CUDA and Pthreads (Multiple GPUs)
 - CUDA and OpenMP
 - CUDA and MPI (GPU cluster)

Programming Models Based on Problem Decomposition

- Data parallelism
 - Splitting data into pieces and letting the processing elements do, in parallel, same task on different pieces
 - Original OpenMP.
- Task Parallelism (MIMD)
 - Decomposing a program into different tasks and execute them on different processing units in parallel.

Data parallelism - Advantages

- Number of parallel processes can be matched to the number of cores
- Can address memory bandwidth bottleneck.
 - Padding and Alignment.
 - Caching
- Safe
 - Synchronization can be handled by the platform runtime rather than the programmer.
 - no deadlocks or race conditions

Data parallelism - Disadvantages

- Some styles of data-parallel processing perform too little work for each memory access and also are limited in expressiveness
- Challenges in the implementation of control flow in some data parallel programming models

Forms of data parallelism

- SIMD
- Vector Processors
- Stream Processing
 - applies a computationally intense kernel to the input and output streams rather than a single operation
 - Stream processing can achieve high arithmetic intensity

Task Parallelism


- Different tasks in parallel
- Usually need to communicate
 - Shared memory
 - Message passing
- Usually need to synchronize
 - Shared memory: locks, barriers, ... etc.
 - Message passing: synchronous and asynchronous
 - Problems: race conditions, deadlocks, scalability problem
- Need to be scheduled
 - Load balancing

Task-Based Programming Models

- A task is a unit of parallel work
- Used in
 - Adopted by OpenMP 3.0
 - Cilk
 - Intel's TBB
 - StarSS, OpenSS, ... etc
- Task dependencies
 - Control-flow model
 - Dataflow model


Task Dependencies Based on Control-Flow

- A task create other task(s) and wait for them to finish
- Suited for recursion and divide-conquer algorithm
- Examples: OpenMP 3.0, Cilk, Intel's TBB



Task Dependencies Based on Dataflow

- Task dependencies is built based on data dependencies
- Task arguments are declared as input, output, or inout
- Task dependency graph
 - directed acyclic graph (DAG)
- More parallelism and less synchronization
- supported SMPs, OMPs, OpenMP 4.0



```
#pragma css task output(a)
void TaskA( float a[M][M]);
```

```
#pragma css task input(a)
void TaskB( float a[M][M]);
```

```
#pragma css task input(a)
void TaskC( float a[M][M]);
```

Task Scheduling


- Goal: minimizes the overall execution time
- Dependencies are represented using a task graph
- Static and dynamic scheduling
 - Assignment to processors at compile time
 - Assignment to processors at runtime
- Load balancing
 - Task pool

OpenMP

- Managed by OpenMP Architecture Review Board (or OpenMP ARB)
- API supports multiplatform using C, C++, and Fortran
- Compiler directives and library routines
- Original OpenMP
 - Data parallelism
- OpenMP 3
 - Support task based approach
 - Dataflow dependency: in|out|inout dependency
- OpenMP 4
 - Support for heterogeneous architectures
 - target construct creates tasks to be executed by accelerators

StarSs and OMPSs

- Source to Source Compiler
- The programmer a serial application with **annotations**
- Dataflow task dependencies
- The runtime system build a data dependency graph



```
#pragma css task output(a)  
void TaskA( float a[M][M]);
```

```
#pragma css task input(a)  
void TaskB( float a[M][M]);
```

```
#pragma css task input(a)  
void TaskC( float a[M][M]);
```

The StarPU project

- LaBRI-INRIA at university of Bordeaux
- Task-based library written in C
- Portability is a goal
 - Dynamic scheduling
 - Several implementations for each task for different architecture
- Schedule of computations on heterogeneous systems
 - Transfer data between CPU and accelerators
 - Determine which tasks to run on CPU and which to run on accelerators based on efficiency

Cilk

- A parallel programming language
- Task based programming model
- Main construct
 - Spawn (fork)
 - Sync (join)
- Control-flow dependencies
- Work-stealing scheduler

Intel's Thread Building Blocks(TBB)

- C++ library
- Task-based
- Constructs for parallel loops, reductions, scans, pipeline parallelism
- Control-flow dependency

Further Parallel Programming Approaches

1. Transactional memory

- Allows a group of load and store instructions to execute in an atomic way
- A scalable alternative to memory locks

2. Resource aware task scheduling

- consider resources like memory and bandwidth

3. Domain specific languages

- Use a *program style that has implicit parallelism*

4. *Pattern Programming Model*

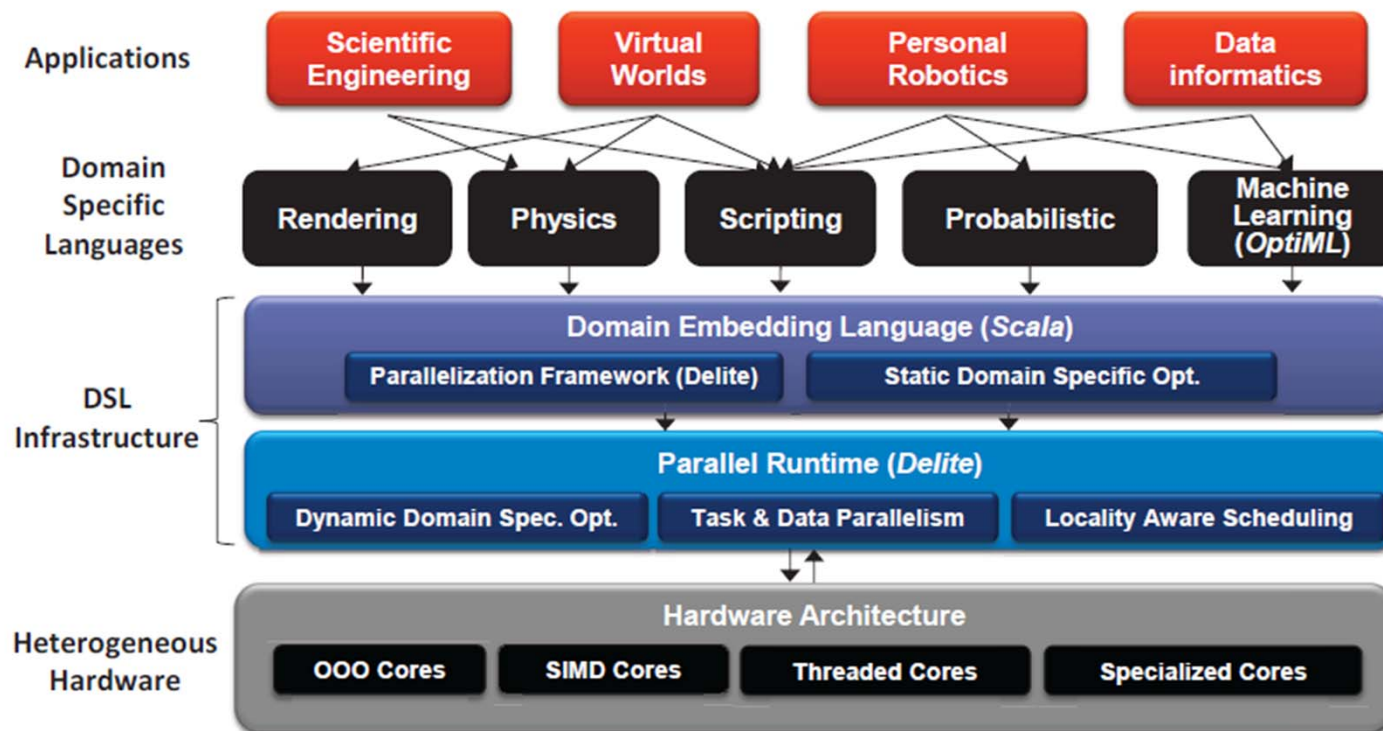
- defined building blocks out of a Pattern Catalogue

Domain-Specific Approach

- A domain-specific language is a computer programming language of **restricted expressiveness** focused on a **particular domain**.
- Examples: TeX, SQL
- Advantages of using the DSL approach
 - support productivity, portability, and performance goals
 - The ability to use domain knowledge to apply static and dynamic optimizations to a program written using a DSL.

Domain-Specific Approach To Heterogeneous Parallelism

Pervasive Parallelism Laboratory at Stanford University



Summary

- The goal of a programming model is to provide a mechanism with which the programmer can specify parallel programs.
- There are many different programming models but also open standard models
- There is a high interest in task-based programming
- Domain-specific approach represents a chance implicit parallelism.

References

- Chafi, H., A. K. Sujeeth, K. J. Brown, H. Lee, A. R. Atreya and K. Olukotun (2011). "A domain-specific approach to heterogeneous parallelism." SIGPLAN Not. **46**(8): 35-46.
- Jahr, R., M. Gerdes and T. Ungerer (2013). A pattern-supported parallelization approach. Proceedings of the 2013 International Workshop on Programming Models and Applications for Multicores and Manycores. Shenzhen, Guangdong, China, ACM: 53-62.
- Javier, D. (2012). "A Survey of Parallel Programming Models and Tools in the Multi and Many-Core Era." IEEE Transactions on Parallel and Distributed Systems 23(8): 1369-1386.
- Project, S. E. (2010). "Parallel Programming Models for Heterogeneous Multicore Architectures." Micro, IEEE **30**(5): 42-53.
- d. Augonnet, S. Thibault, R. Namyst, Pierre-Andr, #x00e9 and Wacrenier (2011). "StarPU: a unified platform for task scheduling on heterogeneous multicore architectures." Concurr. Comput. : Pract. Exper. **23**(2): 187-198.
- Vandierendonck, H., P. Pratikakis and D. S. Nikolopoulos (2011). Parallel programming of general-purpose programs using task-based programming models. Proceedings of the 3rd USENIX conference on Hot topic in parallelism. Berkeley, CA, USENIX Association: 13-13.
- Petrides, P. and P. Trancoso (2013). Addressing the challenges of future large-scale many-core architectures. Proceedings of the ACM International Conference on Computing Frontiers. Ischia, Italy, ACM: 1-4.